Introduction to Machine Learning

Shahane Arushanyan
Machine Learning. Summer 2020

My Background (in Data Science and ML)

- Ucom LLC (2017-2019)
- Philip Morris Armenia (2020-present)

Teaching:

- AUA Machine Learning, Algorithms (2017-2018)
- ISTC (2018)
- ACA (2019-present)

Course Outline

- Intro to Python (optional)
- Numpy, Pandas, Matplotlib
- Regression Algorithms (Linear r., Polynomial r., Ridge r.)
- Classification algorithms (Logistic r., SVM, Decision tree, Random forest)
- Introduction to unsupervised learning
- Introduction to Neural networks

Course Structure

- Lectures
- Practical sessions
- Homeworks (practical + homework)
- 3 Midterm exams (at the end of each month)
- Slack

Lecture Outline

- What is Machine Learning?
- Set ups
- Github

Data Science vs Machine Learning

DS -> using scientific approach to extract meaning and insights from data

ML -> a group of techniques that allow computers to learn from data

What is Machine Learning?

- Data is everywhere and in every field
- Huge amounts of data collected and stored
- Machine learning techniques in our everyday lives

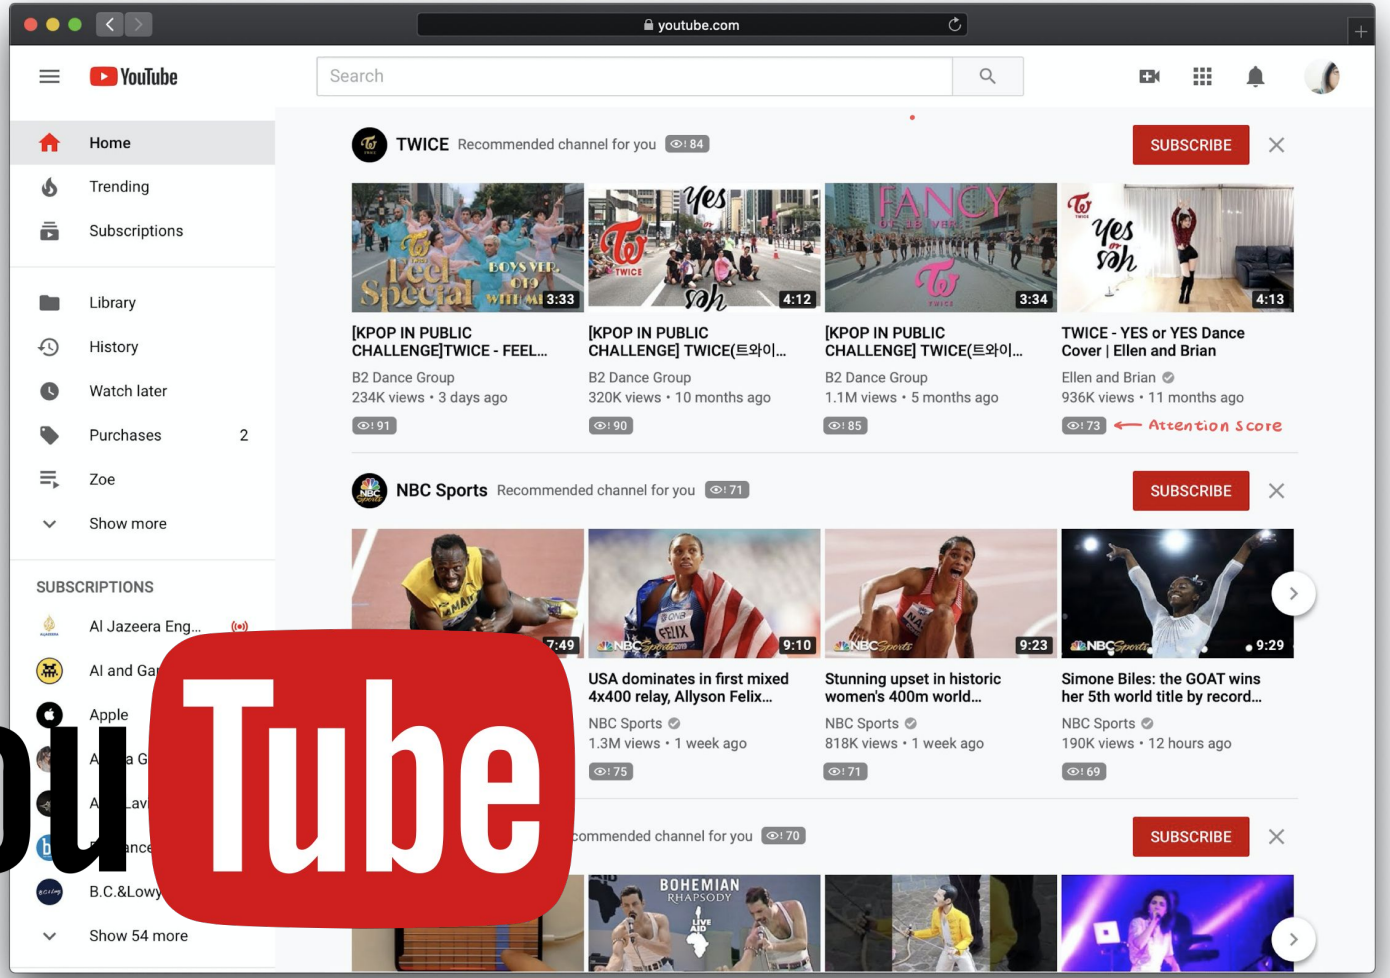
youtube.com
YouTube
Search
Home
Trending
Subscriptions
Library
History
Watch later
Purchases 2
Zoe
Show more
SUBSCRIPTIONS
Al Jazeera Eng...
Apple
B.C.&Lowy
Show 54 more
TWICE Recommended channel for you 84
SUBSCRIBE
[KPOP IN PUBLIC CHALLENGE]TWICE - FEEL...
B2 Dance Group
234K views • 3 days ago
91
3:33
[KPOP IN PUBLIC CHALLENGE] TWICE(트와이...
B2 Dance Group
320K views • 10 months ago
90
4:12
[KPOP IN PUBLIC CHALLENGE] TWICE(트와이...
B2 Dance Group
1.1M views • 5 months ago
85
3:34
TWICE - YES or YES Dance Cover | Ellen and Brian
Ellen and Brian
936K views • 11 months ago
73
Attention Score
4:13
NBC Sports Recommended channel for you 71
SUBSCRIBE
7:49
USA dominates in first mixed 4x400 relay, Allyson Felix...
NBC Sports
1.3M views • 1 week ago
75
9:10
Stunning upset in historic women's 400m world...
NBC Sports
818K views • 1 week ago
71
9:23
Simone Biles: the GOAT wins her 5th world title by record...
NBC Sports
190K views • 12 hours ago
69
9:29
commended channel for you 70
SUBSCRIBE
BOHEMIAN RHAPSODY

YouTube

Face recognition

Vacuum robot

Voice assistant

Movies suggestion

What is Machine Learning?

Machine learning is a study of computer algorithms that improve automatically through experience.

?

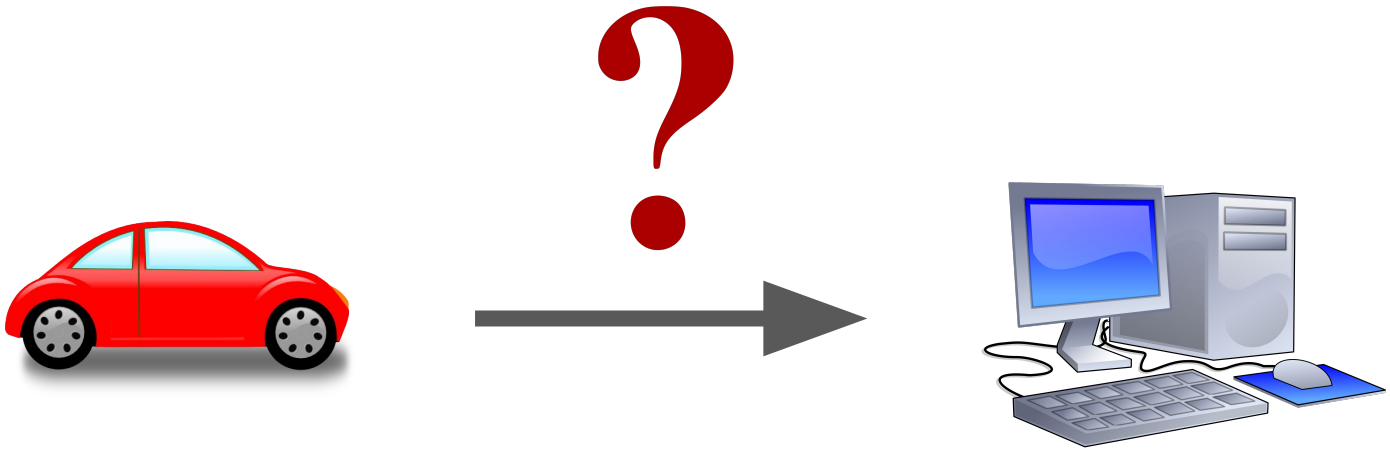

Training Data

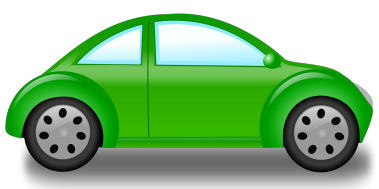

Test Data

$

Size of house

Size of garden

Number of rooms

$ = 1.2 × + 0.7 × + 3.1 ×


$ = A × + B × + C ×
Model with unknown A, B and C to be defined

< Model

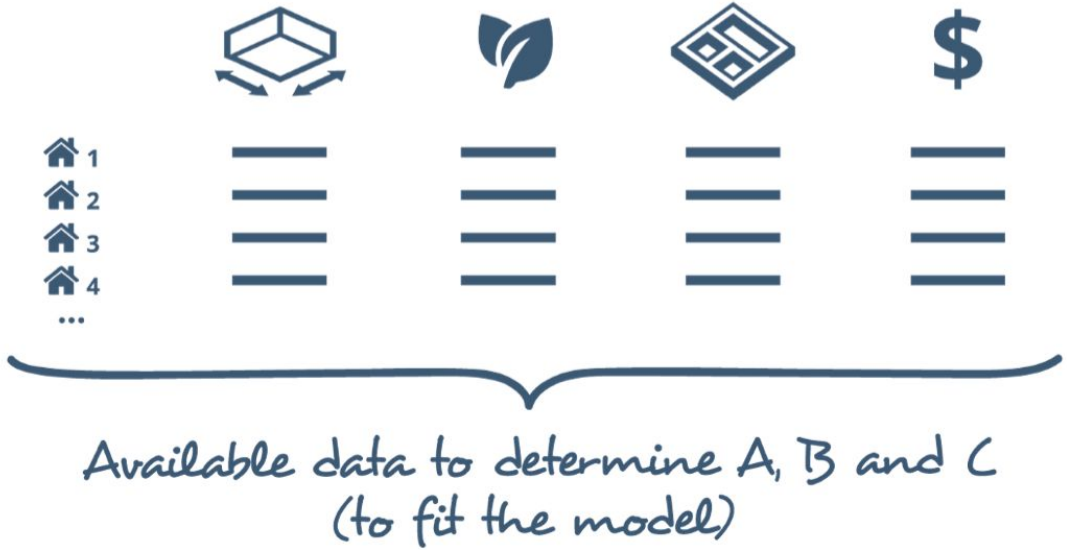
$
1
2
3
4
...
Available data to determine A, B and C
(to fit the model)

< Data

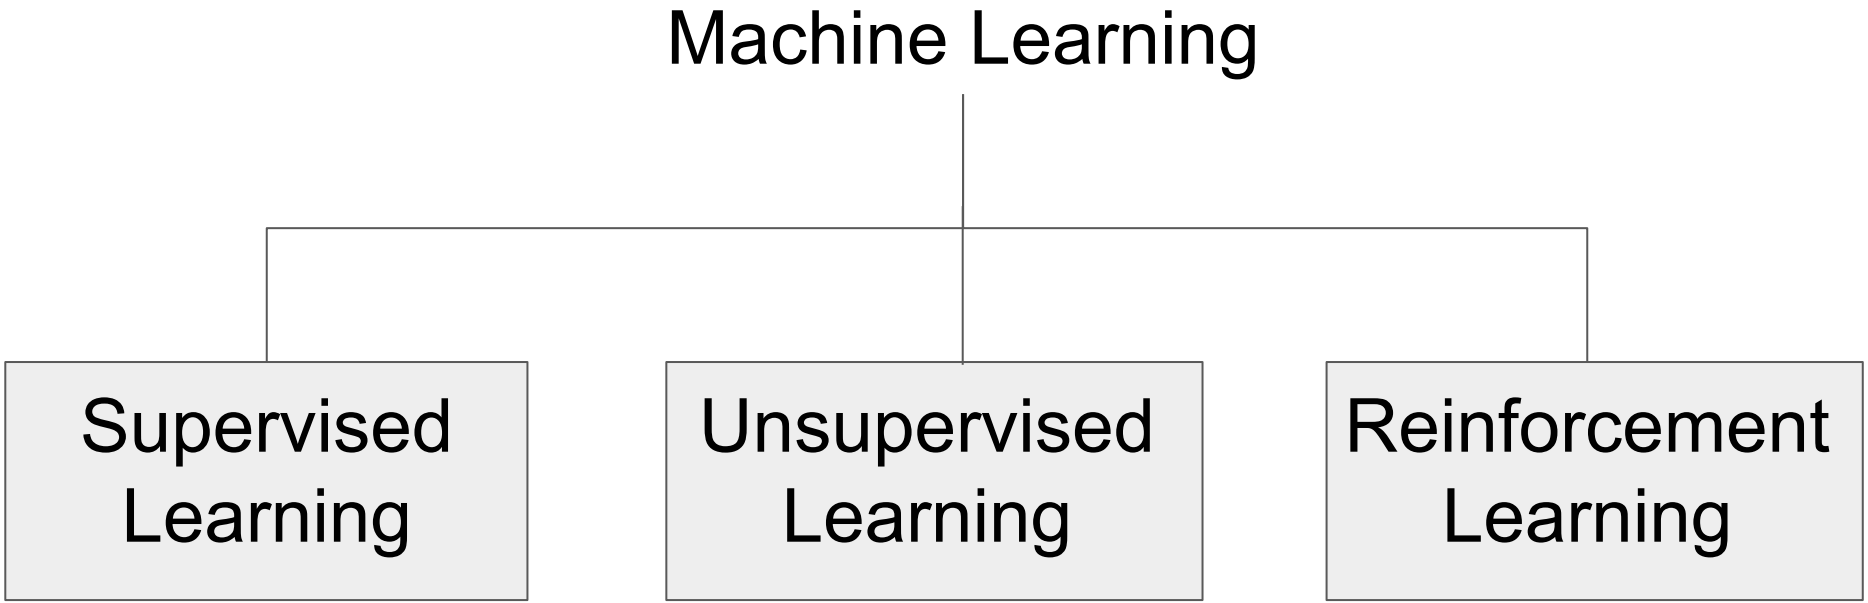
Machine Learning
Supervised Learning
Unsupervised Learning
Reinforcement Learning

Supervised Learning

A supervised model is trained on a labeled dataset of (feature, label) pairs.

Regression Model - numerical label

Problem: Predict weight (number) given height and age

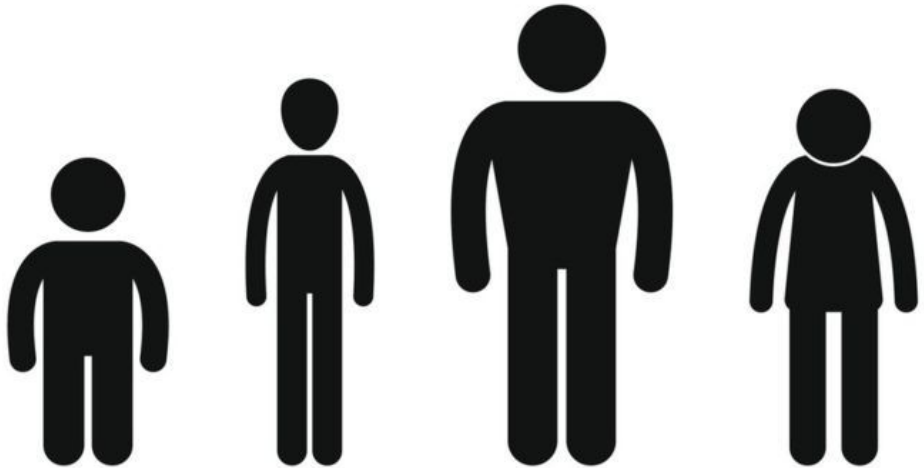

Features:

Height, Age

Label:

Weight

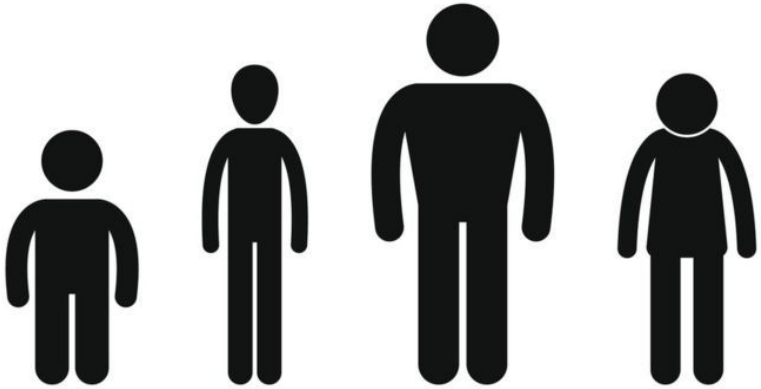

Height:	1.50	1.70	2.10	1.55	1.62
Age:	10	24	40	20	30
Weight:	40	58	80	45	**?**

Training data

Test data

Classification Model - categorical label

Problem: Predict if the object is an apple or not (True/False) given color and shape.

Features:

Color, Shape

Label:

True/False

Color:	red	yellow	green	yellow	Green
Shape:	round	round	oval	oval	round
Apple:	True	True	False	False	**?**

Training data

Test data

Supervised Learning

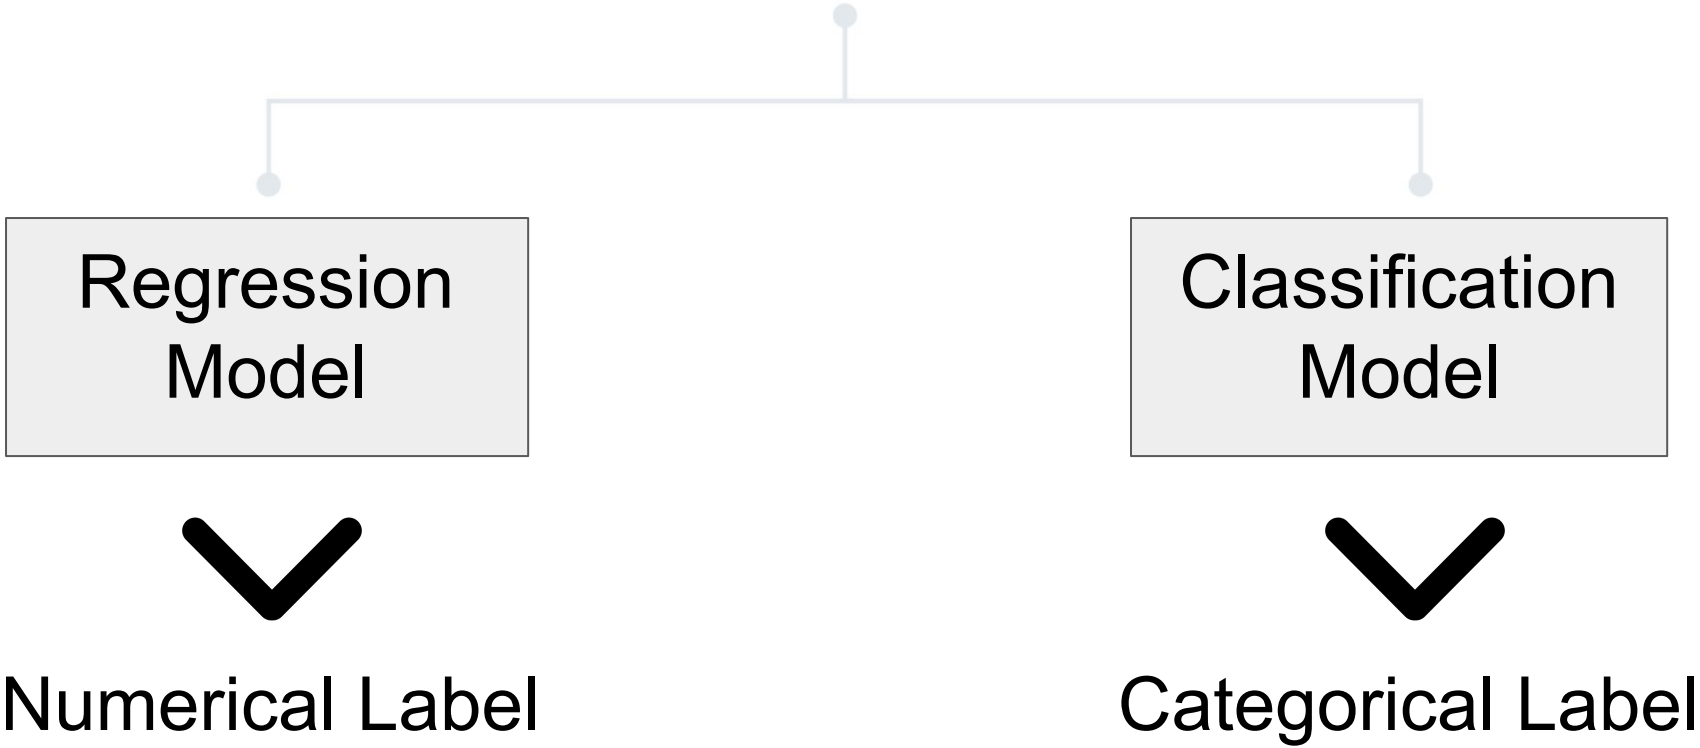

Unsupervised Learning

An unsupervised model is trained on a unlabeled dataset that contains only features but with NO labels

Clustering model - group similar instances together

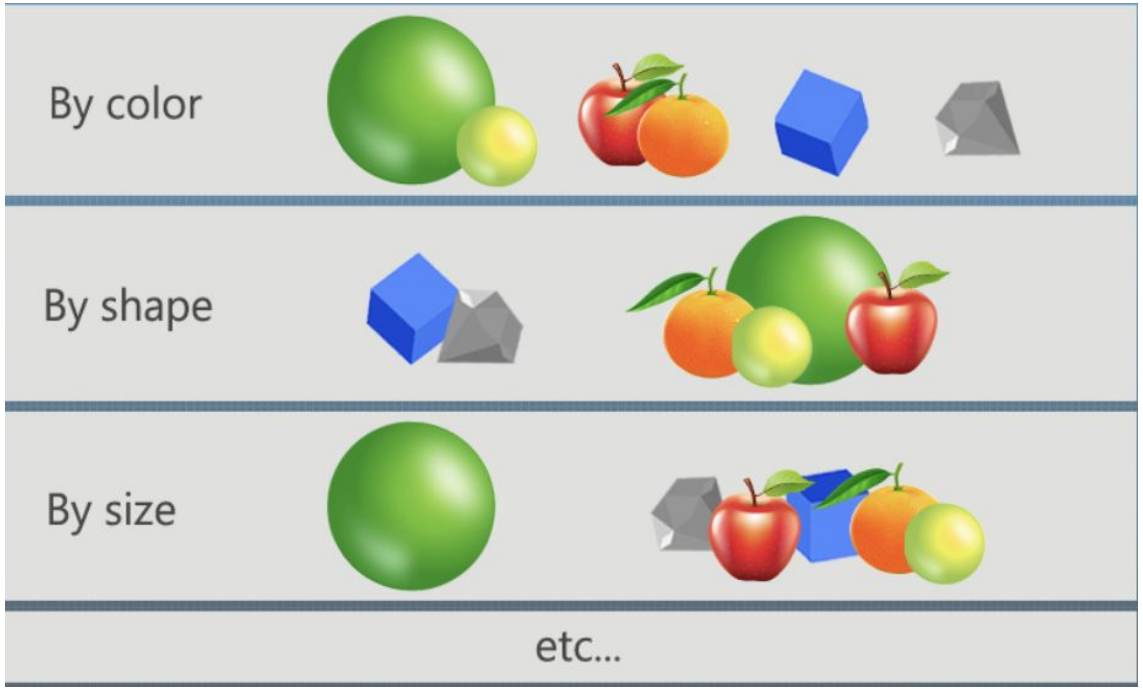

Clustering model - customer segmentation given income and age

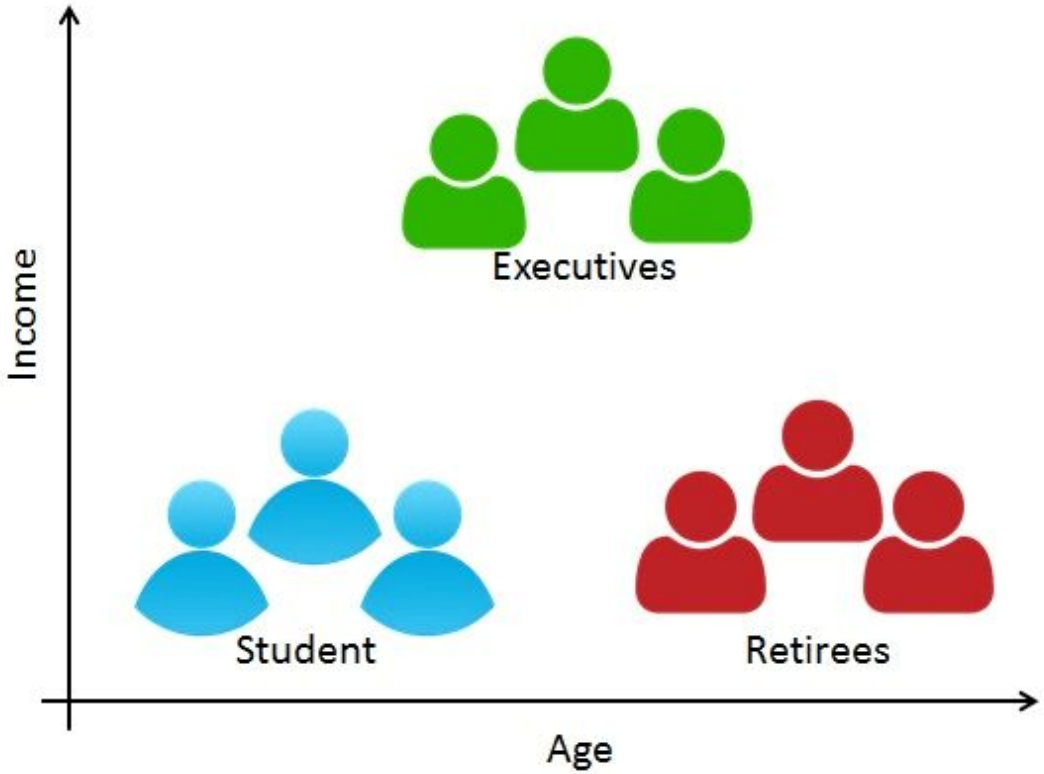

Dimension reduction model - express data with 2-3 dimensions

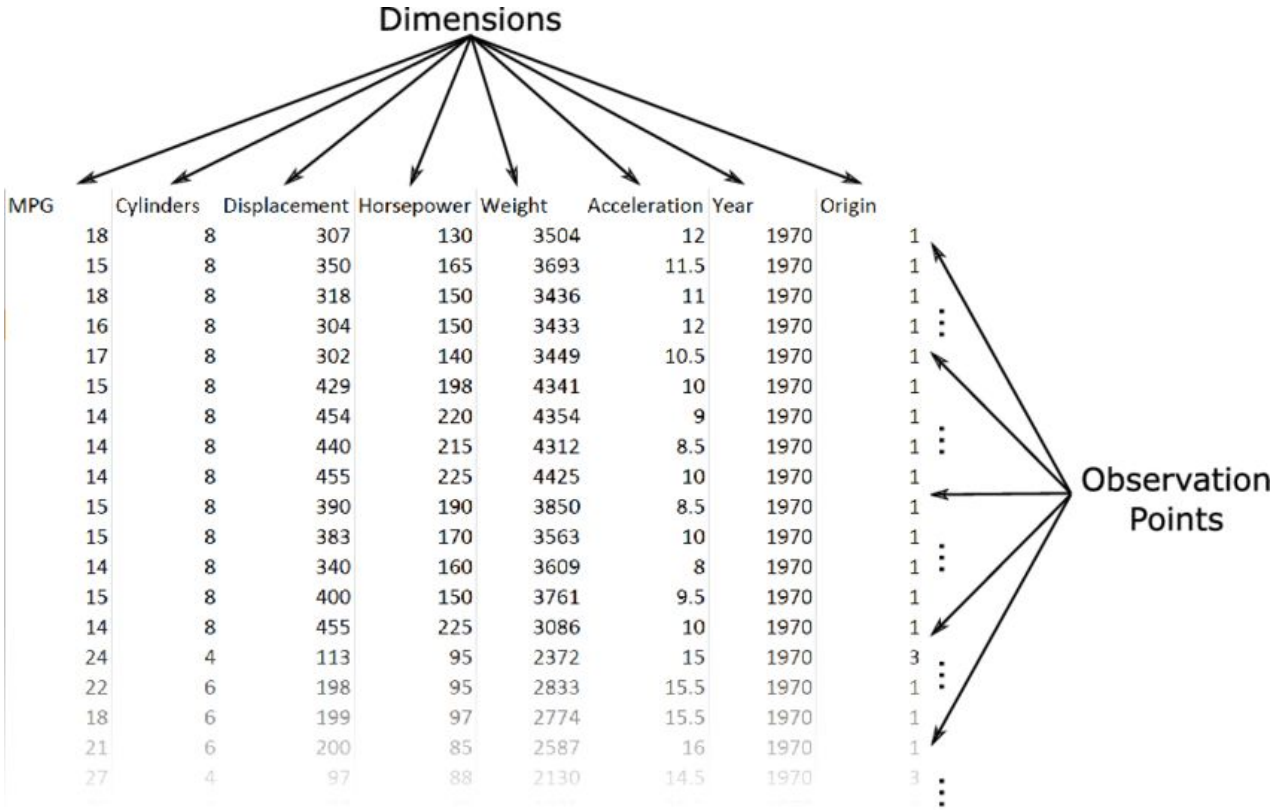

Dimension reduction model - express data with 2-3 dimensions

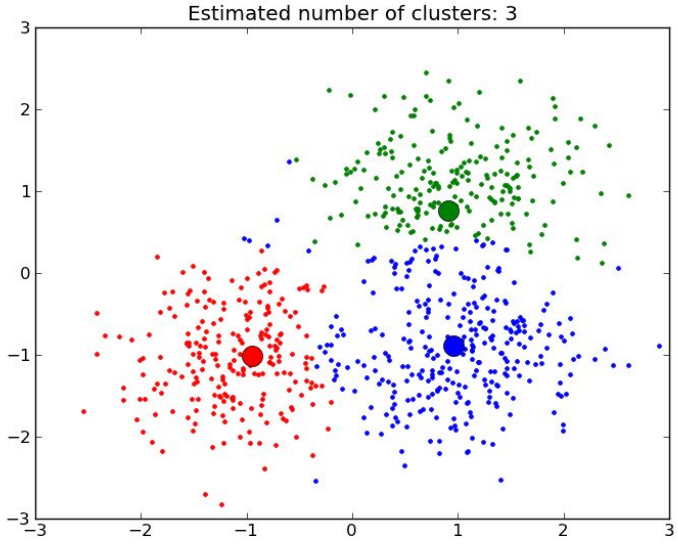

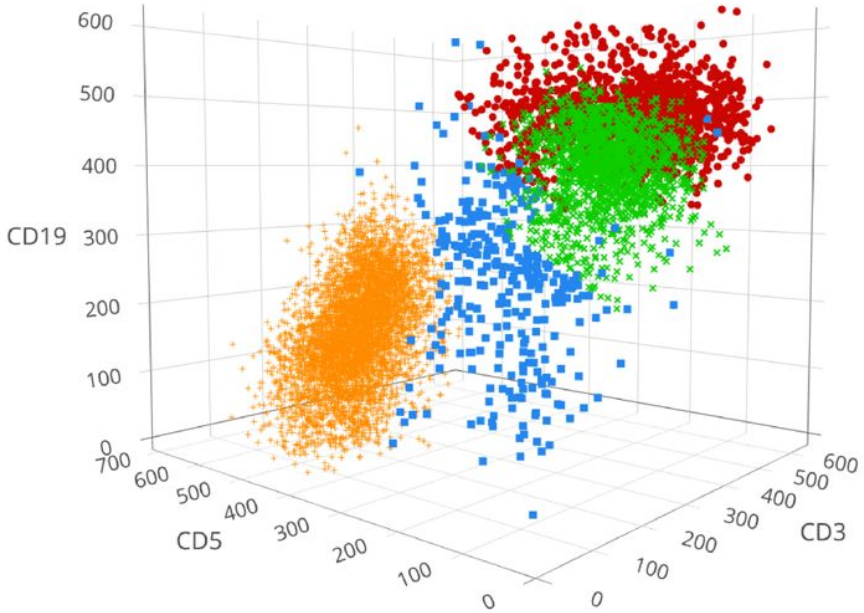

Anomaly detection - finding outliers

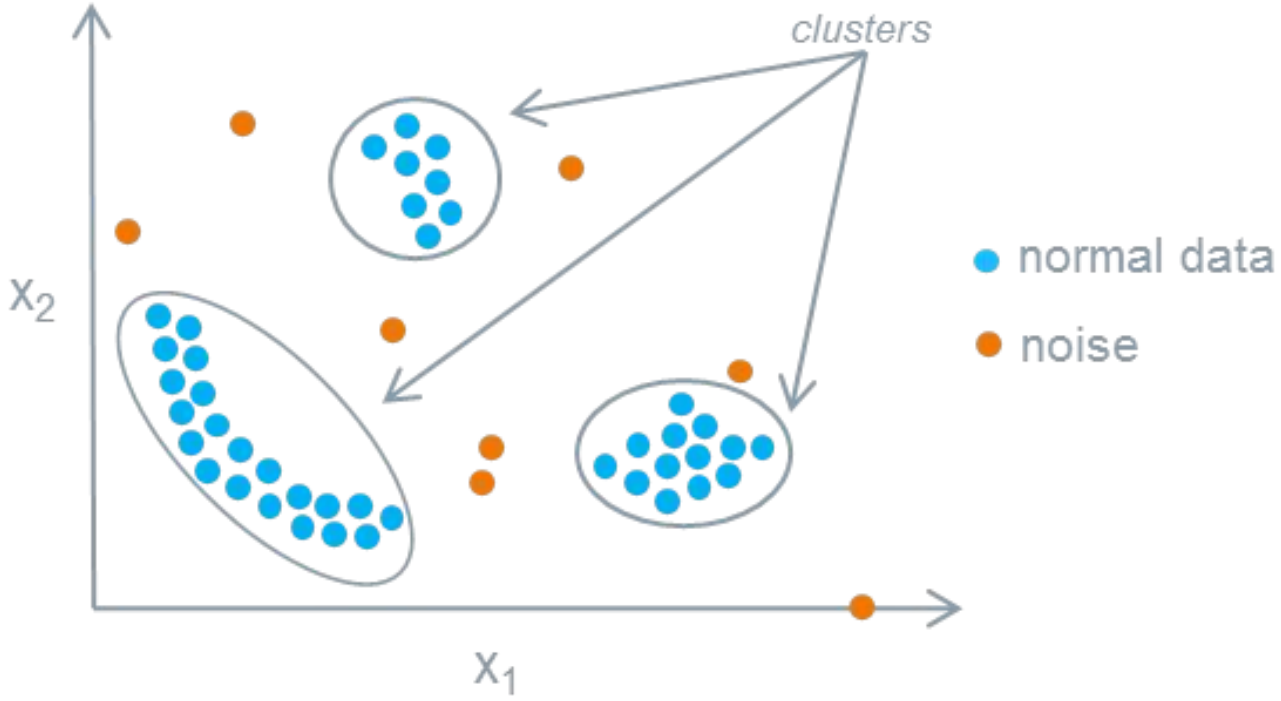

Anomaly detection - finding outliers

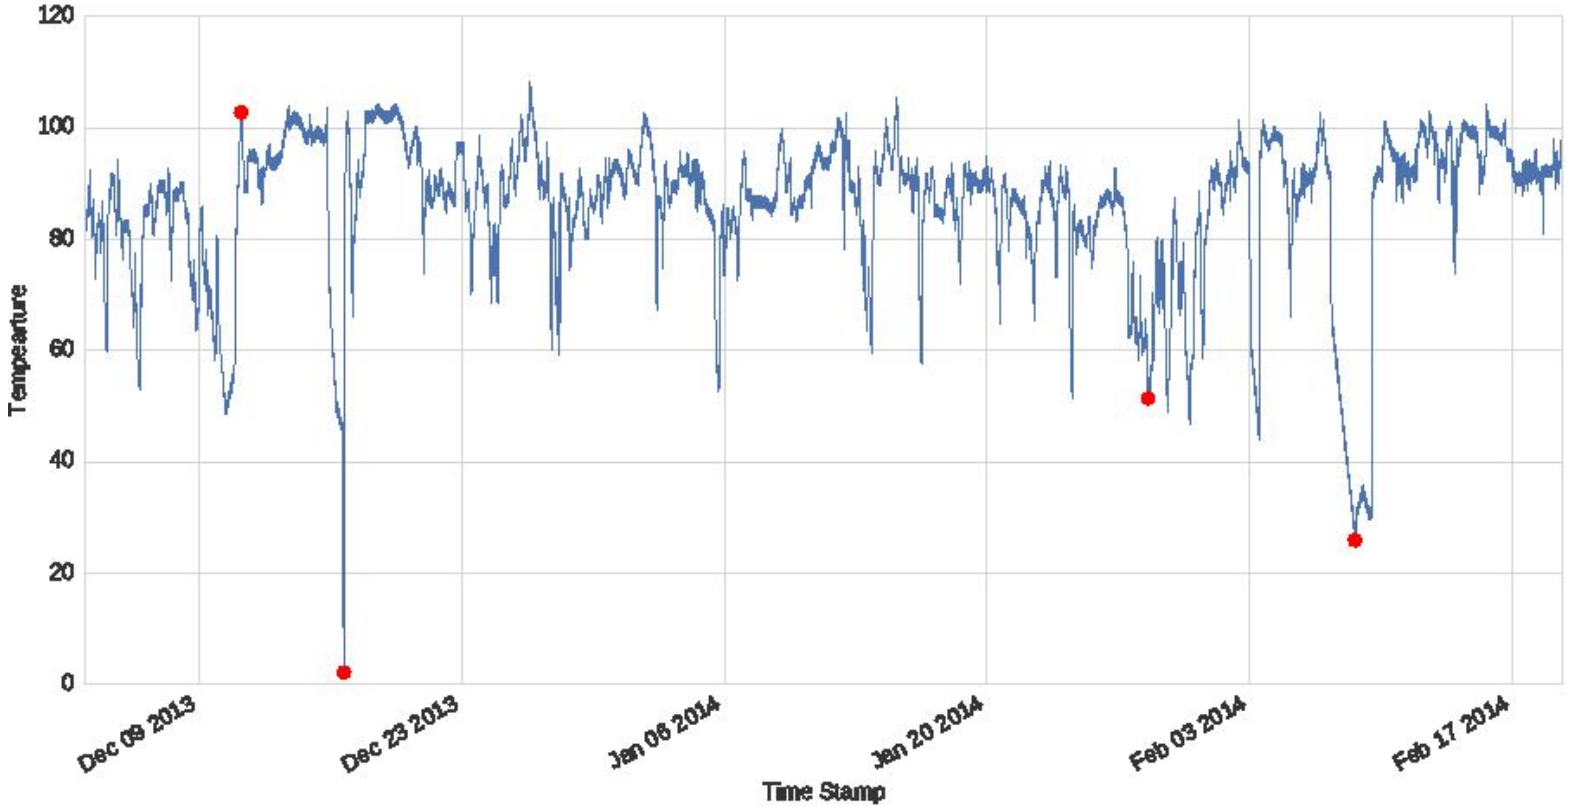

Neural Networks

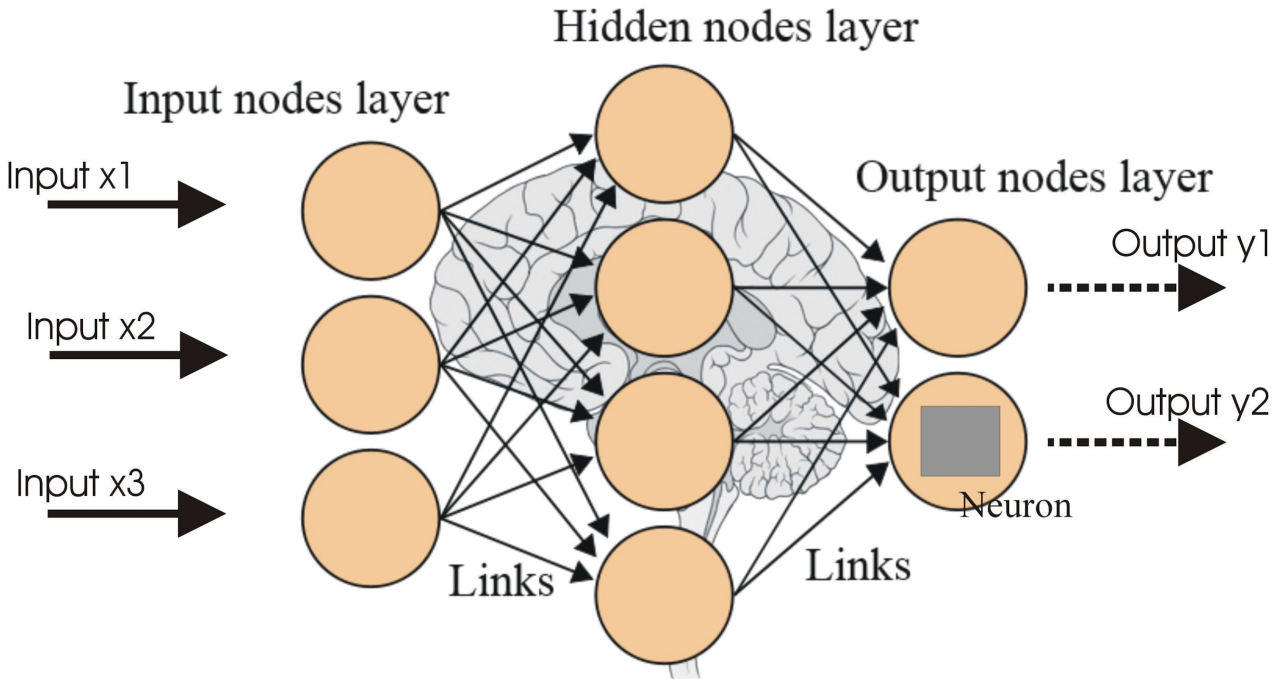

Convolutional Neural Networks

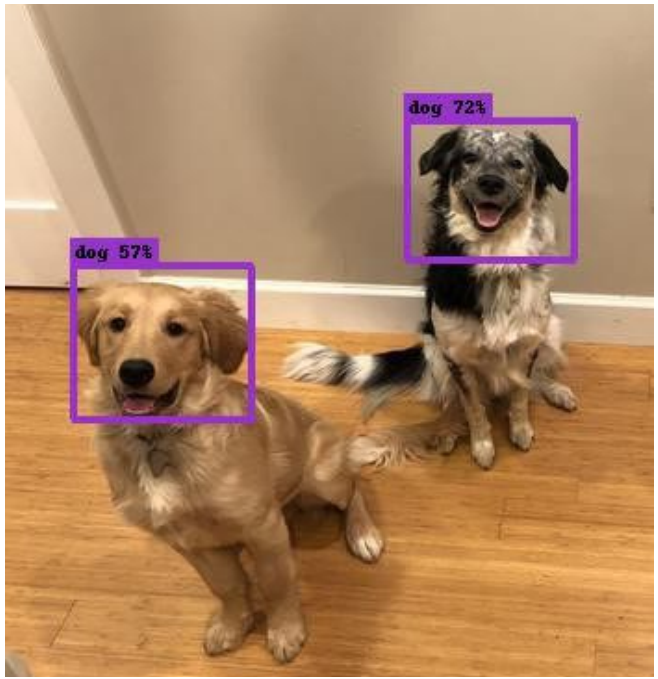

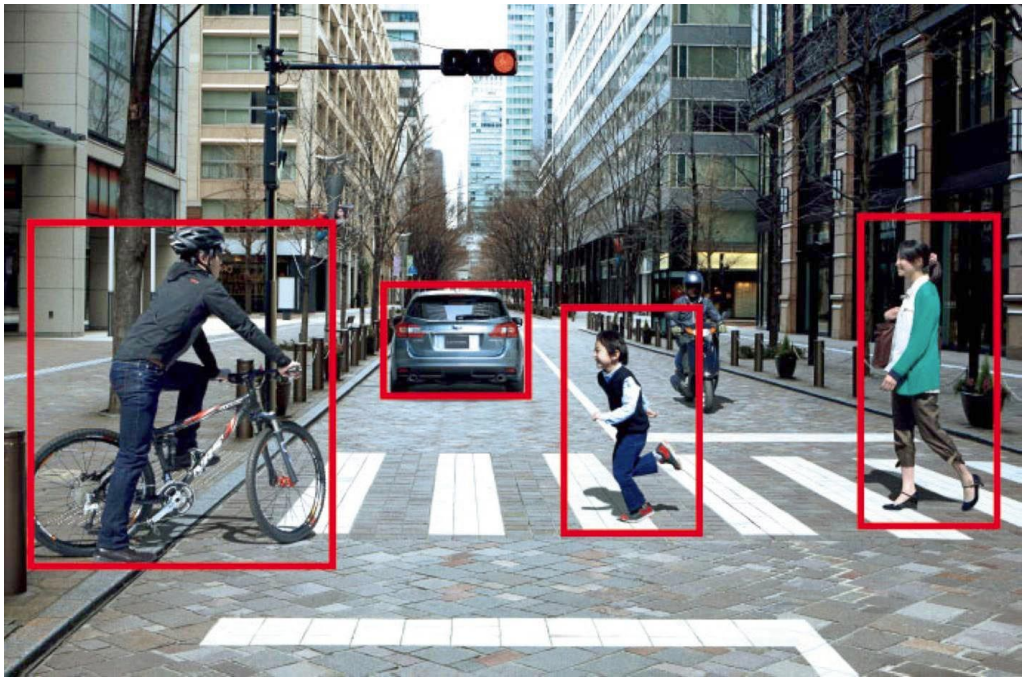

Reinforcement Learning

Train a machine learning model to generate a sequence of decisions

(Model)

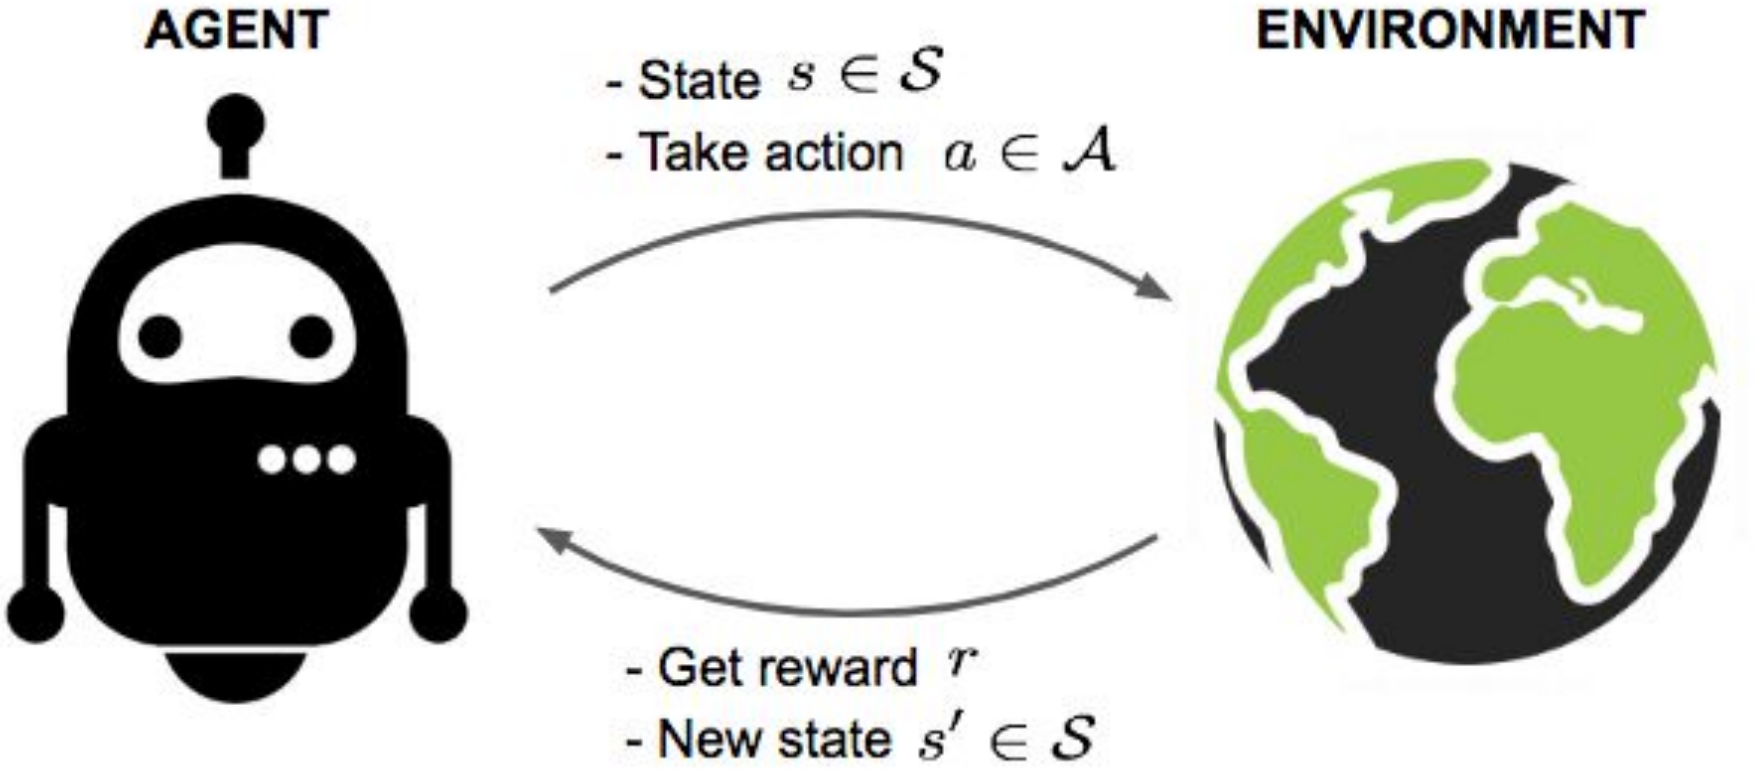
AGENT
ENVIRONMENT
- State $s \in \mathcal{S}$
- Take action $a \in \mathcal{A}$
- Get reward r
- New state $s' \in \mathcal{S}$

Mario game

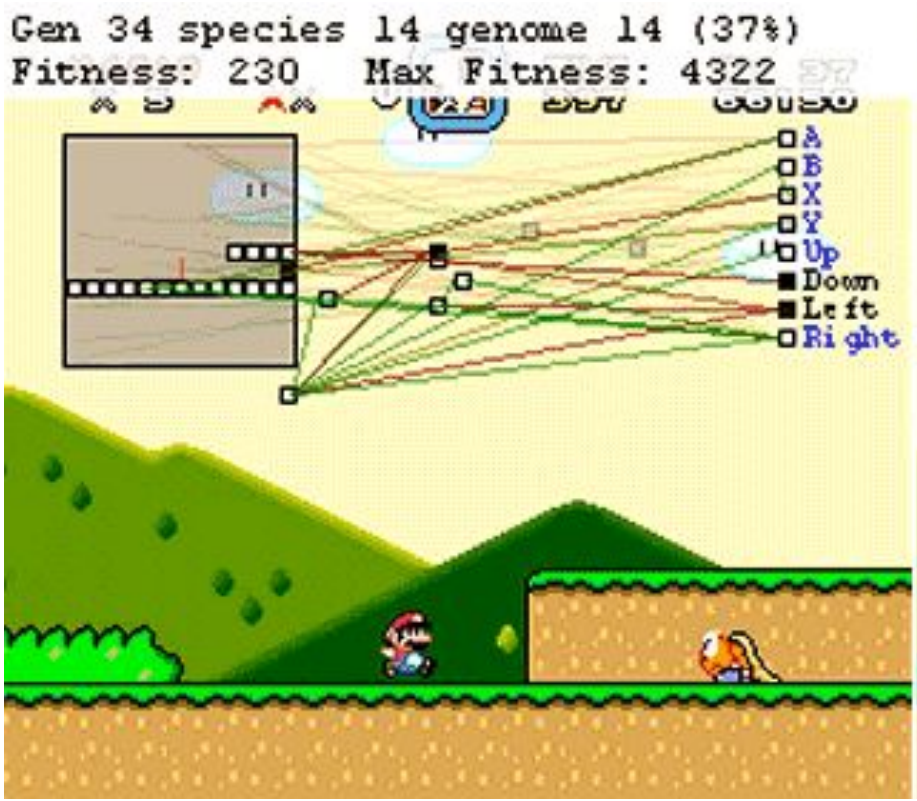
Gen 34 species 14 genome 14 (37%)
Fitness: 230 Max Fitness: 4322
A
B
X
Y
Up
Down
Left
Right

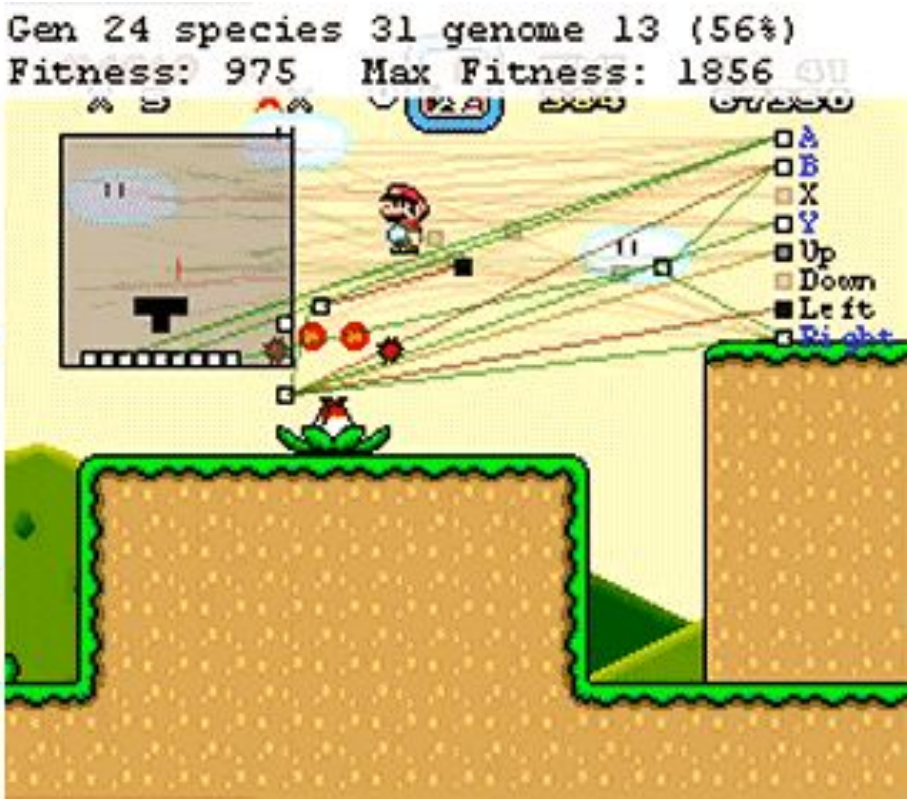
Gen 24 species 31 genome 13 (56%)
Fitness: 975 Max Fitness: 1856
A
B
X
Y
Up
Down
Left

Autonomous car navigation

Environment: street model

Set of actions: 

Scoring: penalty/reward

https://www.youtube.com/watch?v=3ROVzjkkCIA

Thank you!